Multilabel Classification of Research Articles

Teo Zhan Rui

29th April 2021

Problem Statement

1) What are the challenges of multilabel classification?

2) What are the suitable metrics for measuring performance?

3) Which model produced the best scores and are suitable for deployment?

Workflow

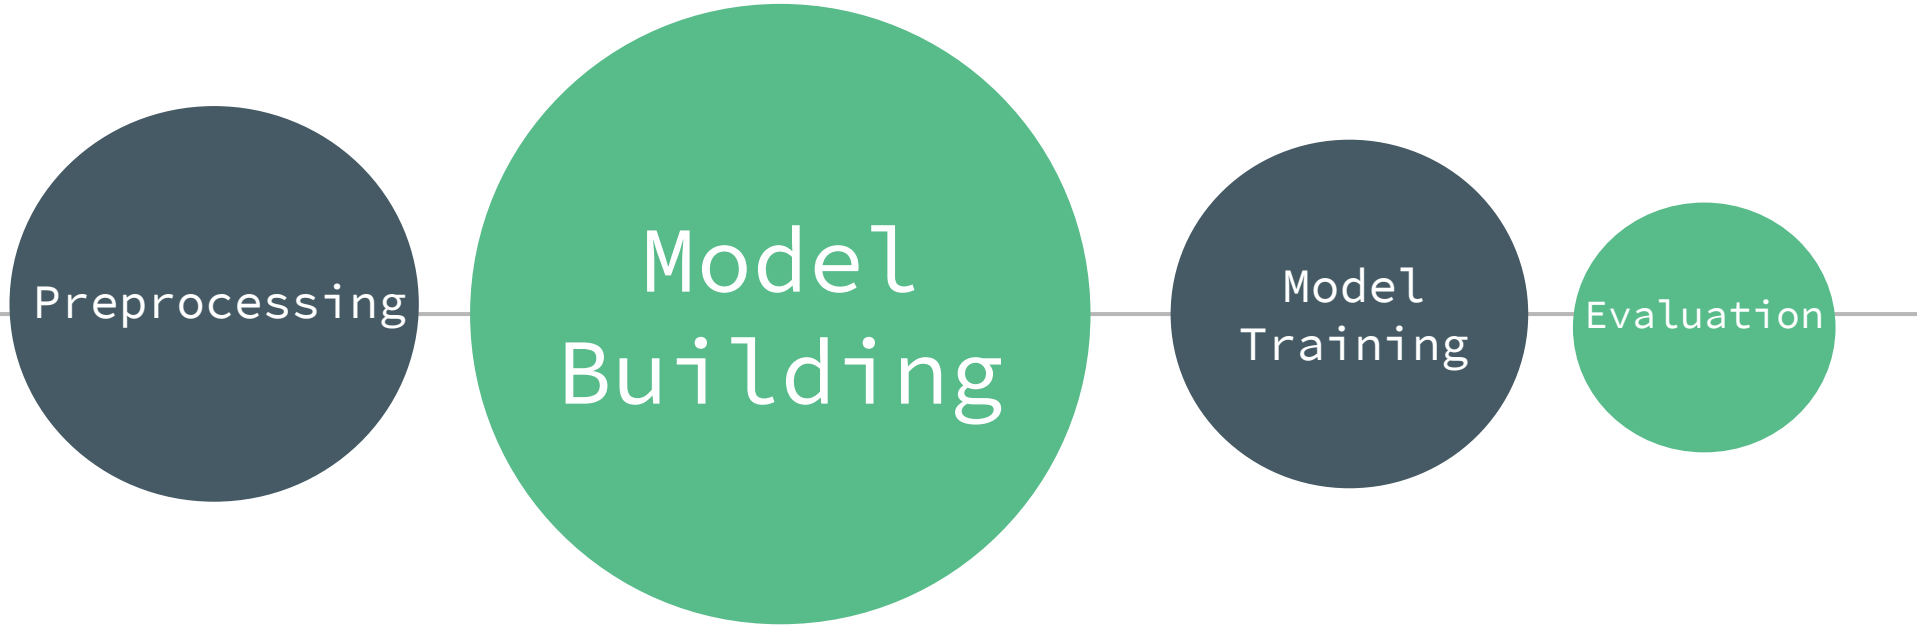

Data

number of articles
20792

maximum length of text
492

total number of labels
6

What is Multilabel Classification?

	Com Sci	Physics	Math	Statistics	Quant. Biology	Quant. Finance
article 1	**1**	0	**1**	**1**	0	0
article 2	0	**1**	**1**	0	0	0
article 3	0	0	0	0	**1**	0

Which metrics to use?

$$\textbf{micro F1} = \frac{2 \times \text{precision} \times \text{recall}}{\text{precision} + \text{recall}}$$

$$\text{precision} = \frac{TP}{TP + FP}$$

$$\text{recall} = \frac{TP}{TP + FN}$$

Sum up TP, FP and FN for each class.

Which metrics to use?

	Com Sci	Physics	Math
ground truth	**1**	**0**	**1**
prediction	**0**	**1**	**1**

hamming loss = 2/3

all labels correct :

hamming loss = 0/3

all labels wrong :

hamming loss = 3/3

How many labels do the articles have?

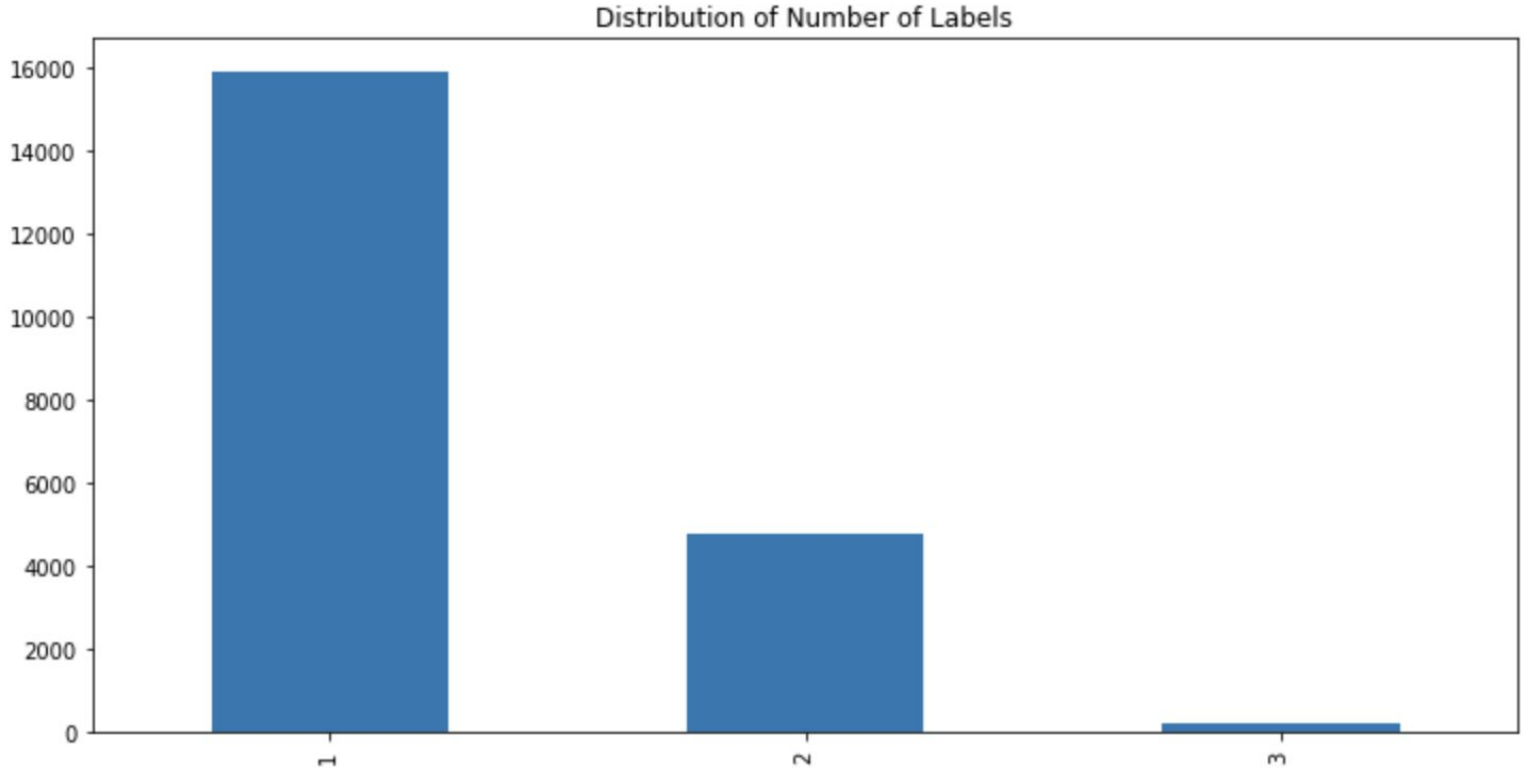

How many observations for each class?

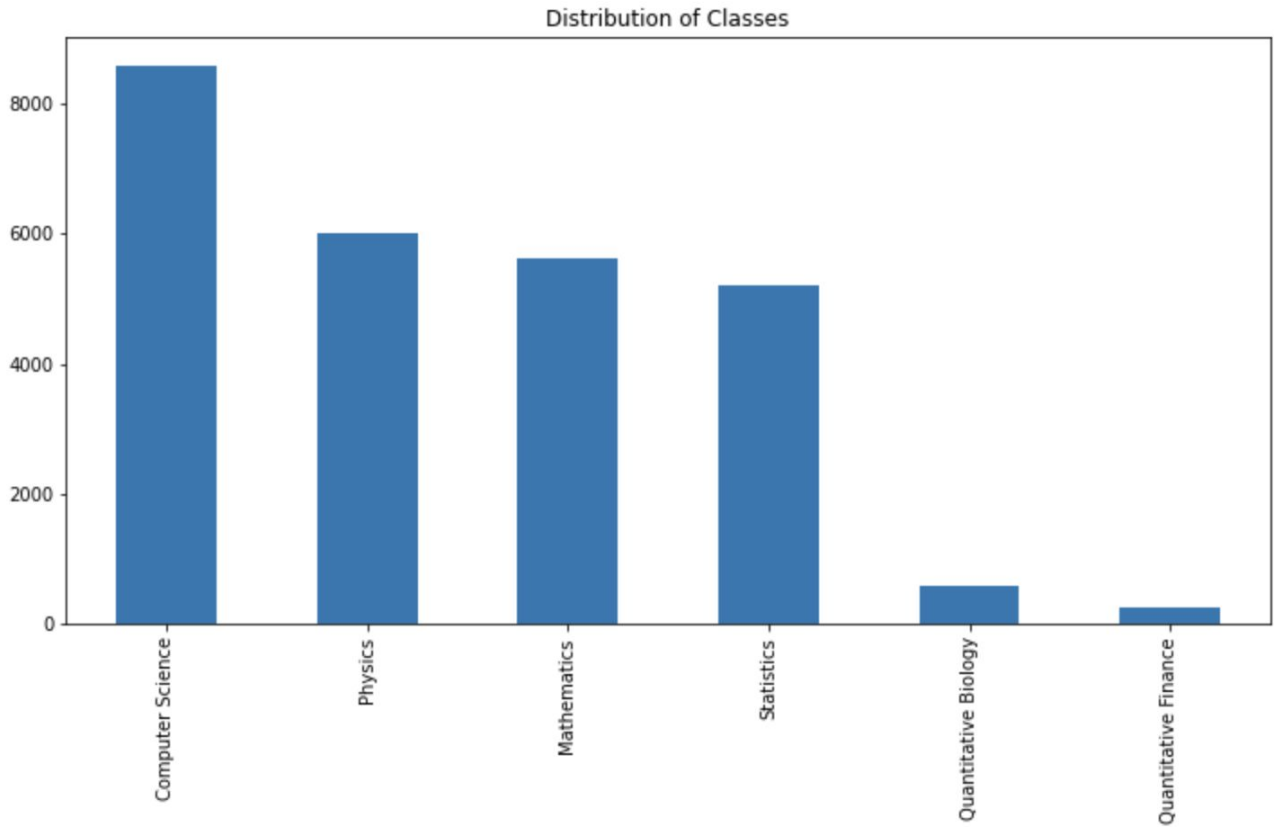

Visualizing in 2D

Transform TFIDF into 2 dimension vectors using PCA

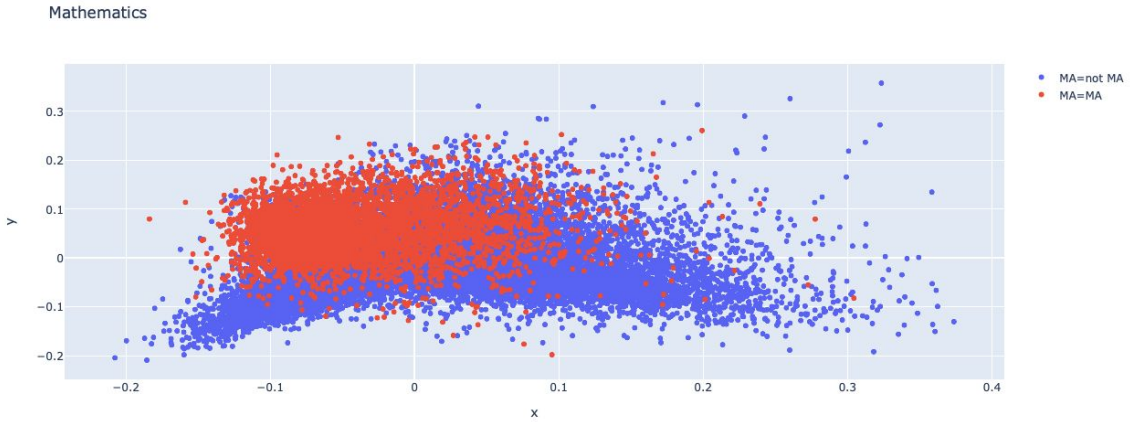

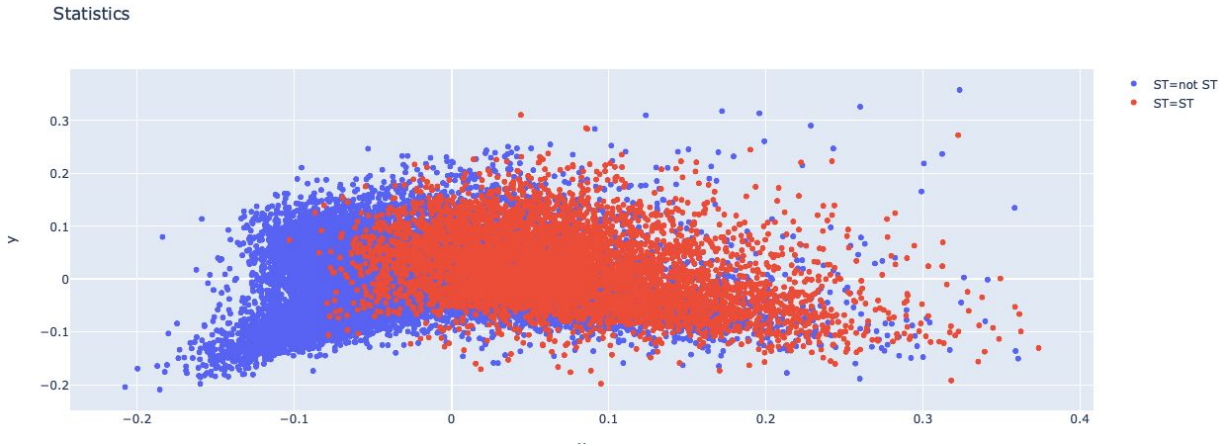

Visualizing in 2D

Transform TFIDF into 2 dimension vectors using PCA

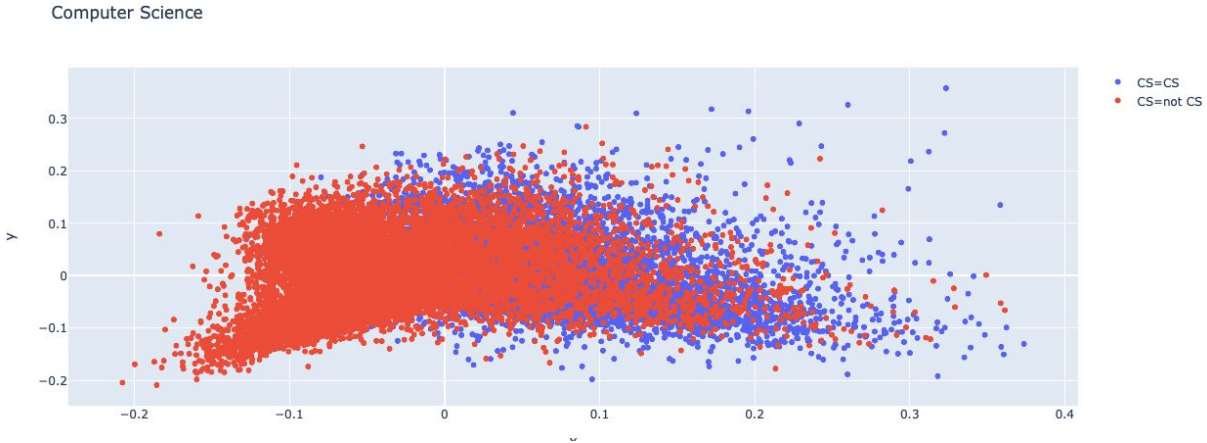

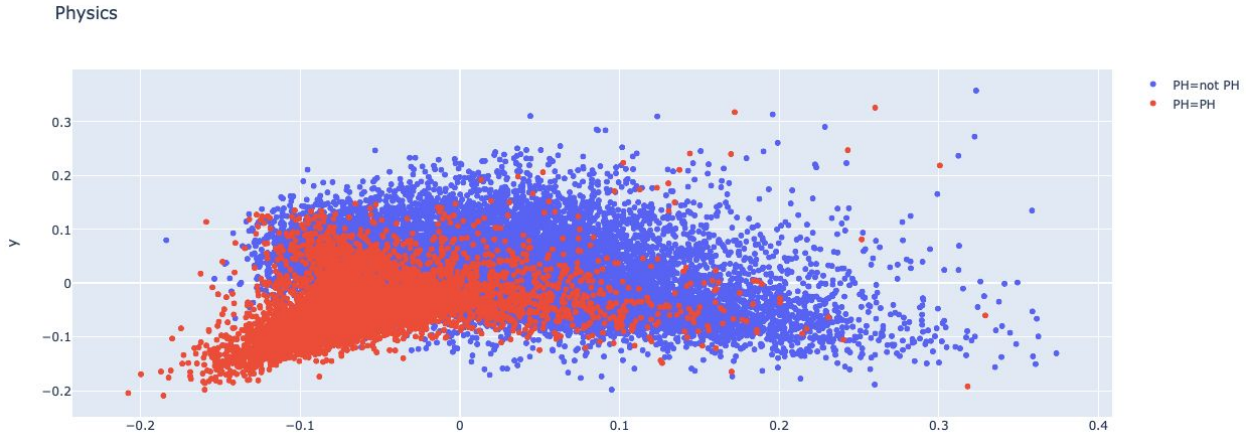

Visualizing in 2D

Transform TFIDF into 2 dimension vectors using PCA

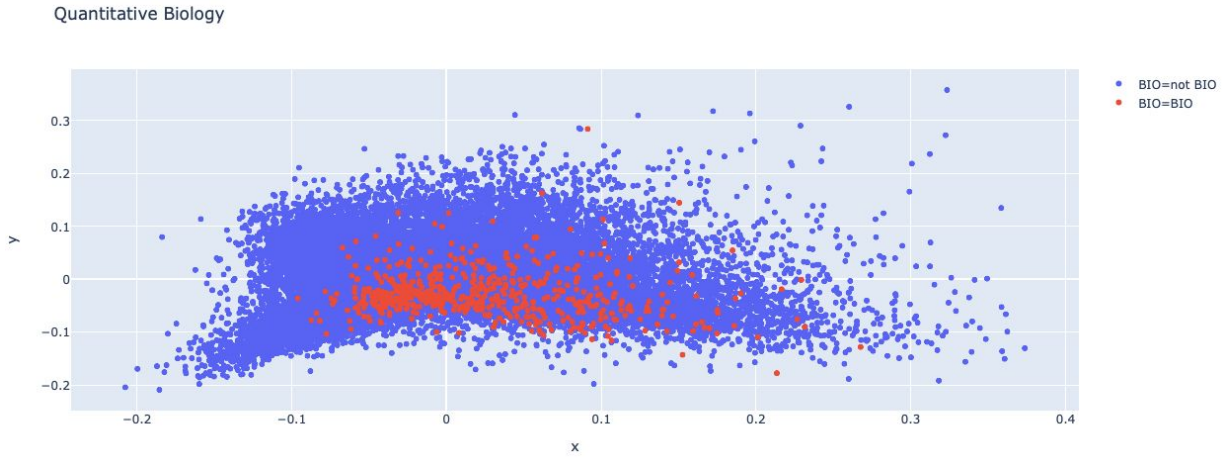

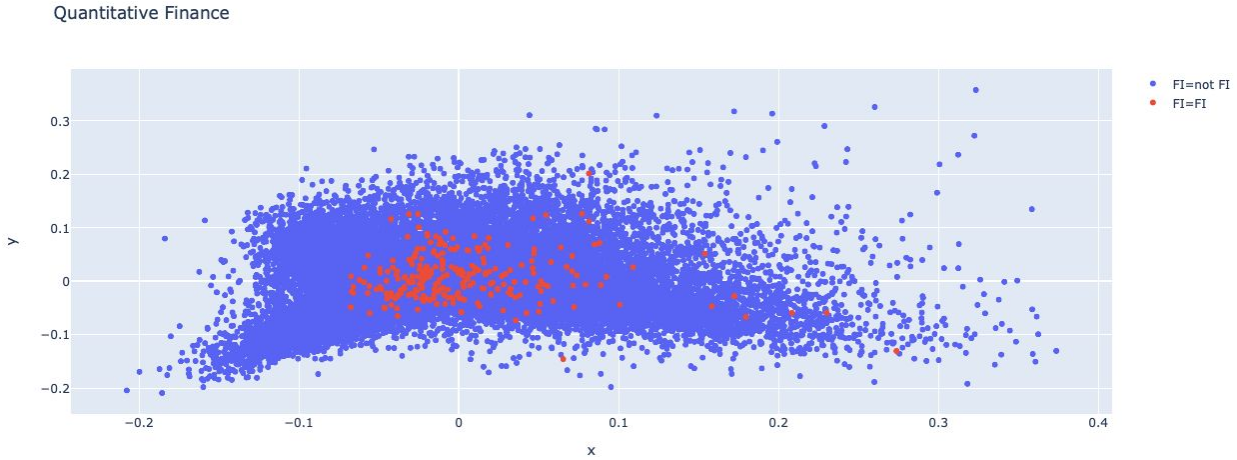

Sklearn

Logistic Regression - Benchmark

	precision	recall	f1-score	support
Computer Science	0.822	0.830	0.826	1692
Physics	0.936	0.803	0.865	1226
Mathematics	0.874	0.733	0.798	1150
Statistics	0.828	0.664	0.737	1069
Quantitative Biology	0.625	0.041	0.077	122
Quantitative Finance	1.000	0.200	0.333	45
micro avg	0.861	0.746	0.799	5304
macro avg	0.848	0.545	0.606	5304
weighted avg	0.858	0.746	0.790	5304
samples avg	0.807	0.781	0.778	5304

hamming loss : 0.07902264600715136

Logistic Regression

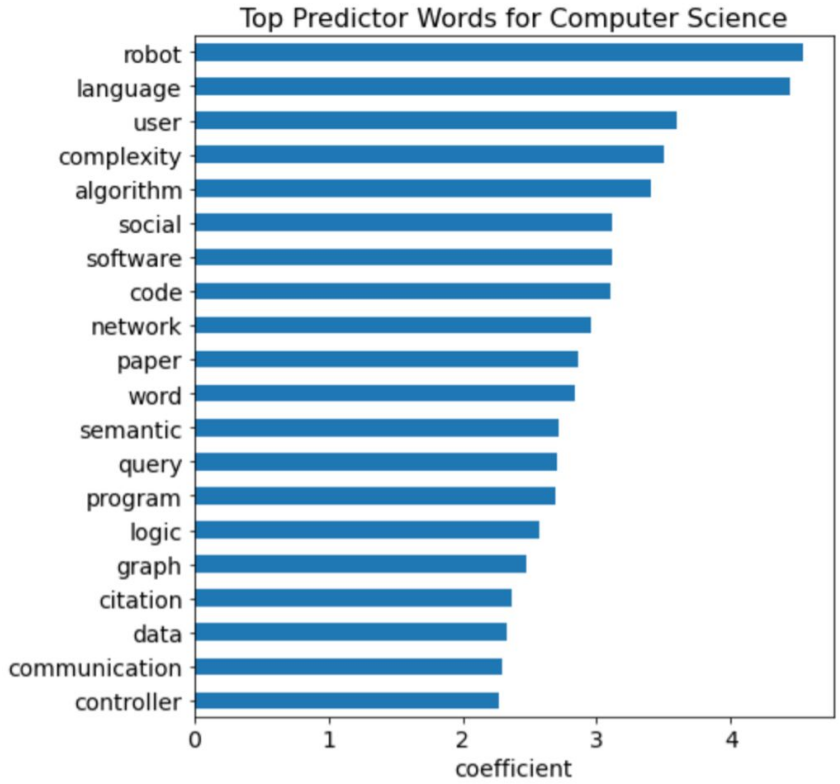

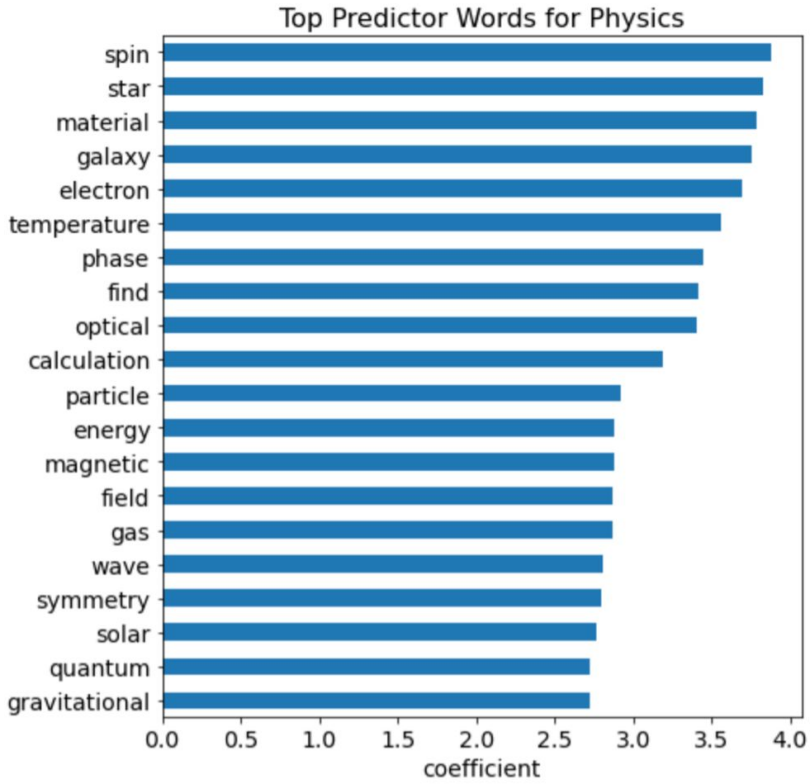

Logistic Regression

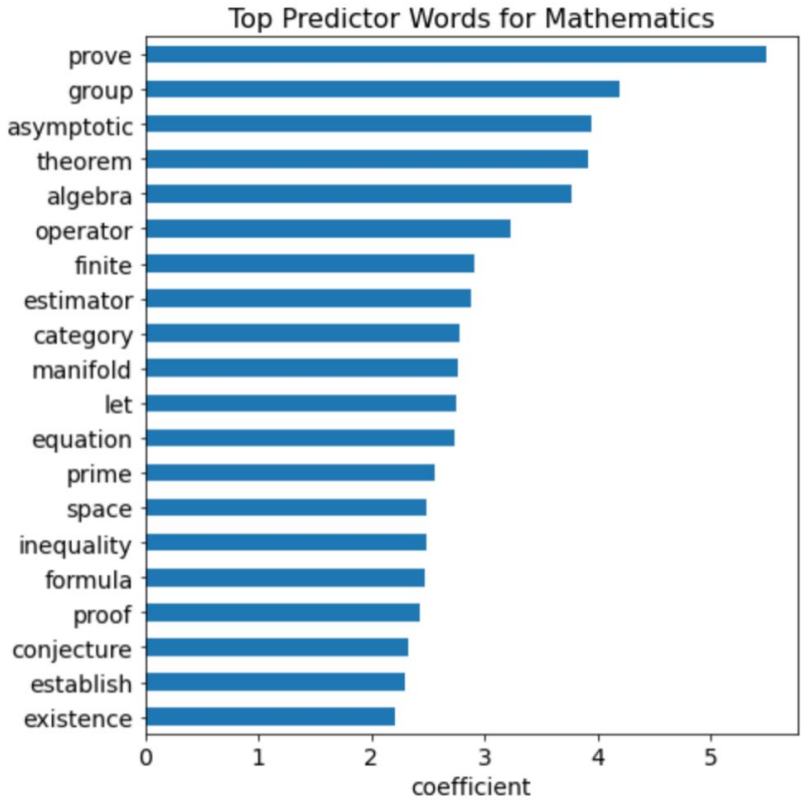

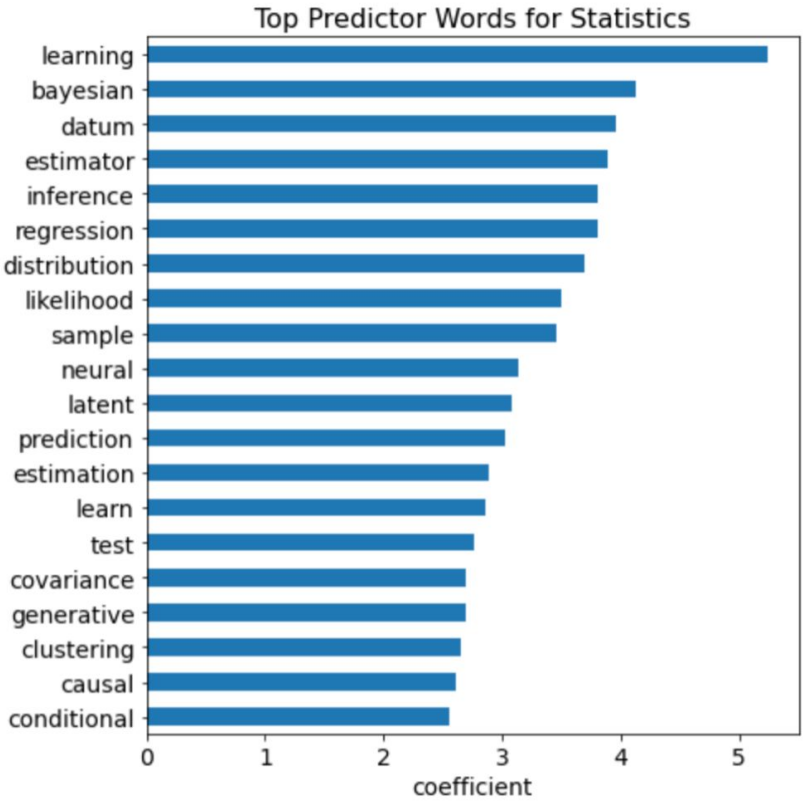

Logistic Regression

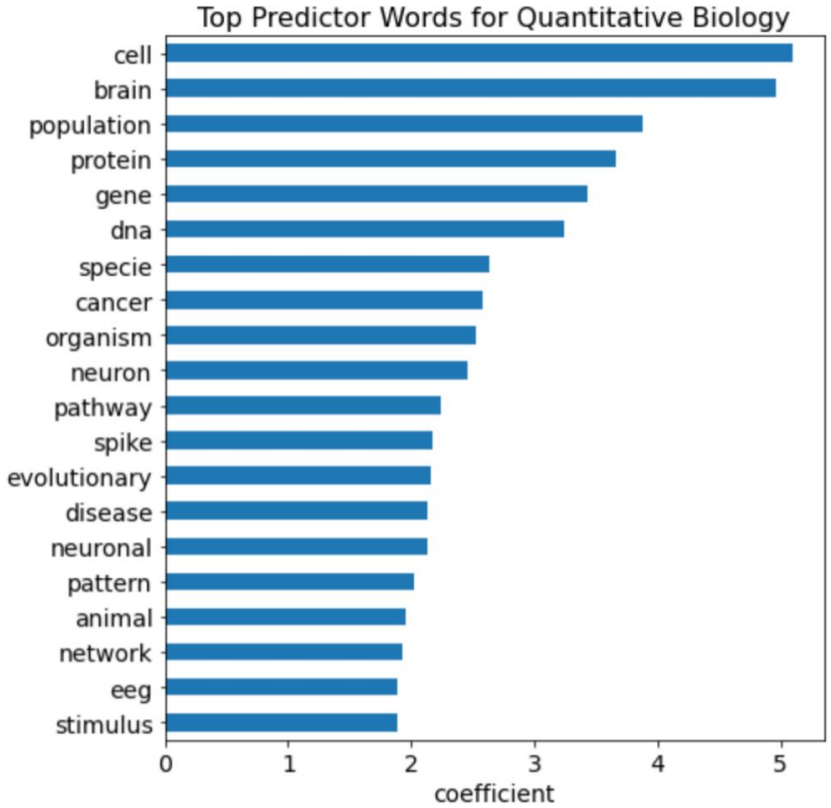

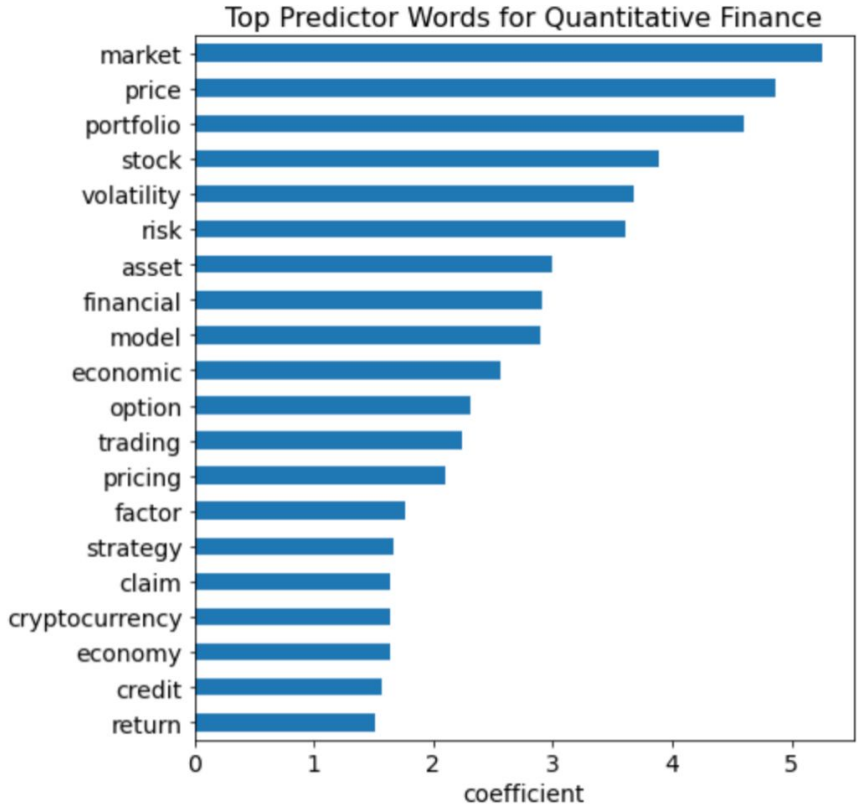

Are sklearn models suitable?

1) does not support multilabel
2) One vs Rest classifier - 6 different models
3) SVM - 2 minutes to complete prediction on 4000 validation articles (X_test)

	Com Sci	Physics	Math	Statistics	Quant. Biology	Quant. Finance
article 1	1	0	1	1	0	0
article 2	0	1	1	0	0	0
article 3	0	0	0	0	1	0

Recurrent Neural Networks

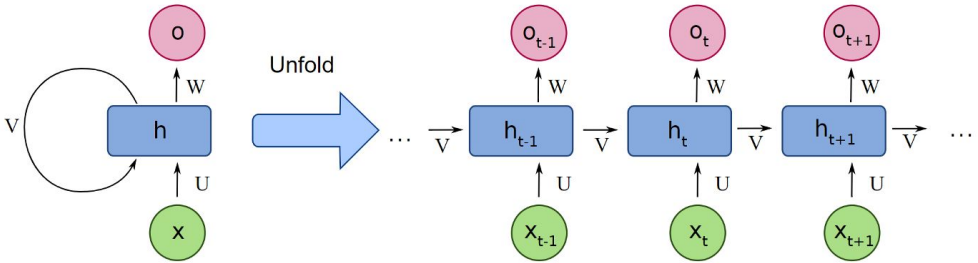

Source:
https://medium.com/deeplearningbrasilia/deep-learning-recurrent-neural-networks-f9482a24d010

Word2Vec and LexVec

- representation of words in 300 dimension vectors
- compare similar words
- self training vs pre-trained

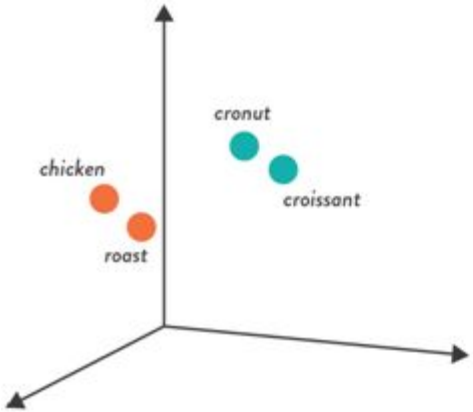

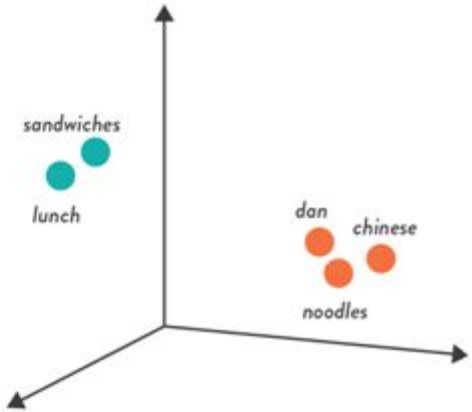

Source: https://www.gsitechnology.com/Leveraging-Word2vec-for-More-than-Text

Word2Vec and LexVec

‘bayes’ :

CBOW	Skipgram	LexVec
dirichlet distribution	gp sum	bayesian
selection procedure	sure convergence	regression
posterior sampling	hellinger distance	inference
sure convergence	entropy sgd	probabilistic
frequentist	likelihood bootstrap	probability

Word2Vec and LexVec

‘cat’ :

CBOW	Skipgram	LexVec
dog	defence	dog
traffic sign	debugger	cats
robot assist	curvature bound	feline
large annotate	dog	puppy
handwritten character	fully autonomous	kitten

My Neural Networks - Multilabel

- output layer : 6
- activation : sigmoid
- optimizer : adam
- loss : binary crossentropy

	Com Sci	Physics	Math	Statistics	Quant. Biology	Quant. Finance
article 1	1	0	1	1	0	0
article 2	0	1	1	0	0	0
article 3	0	0	0	0	1	0

RNN GRU LSTM

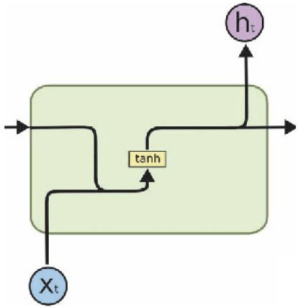

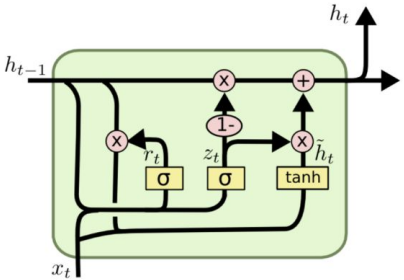

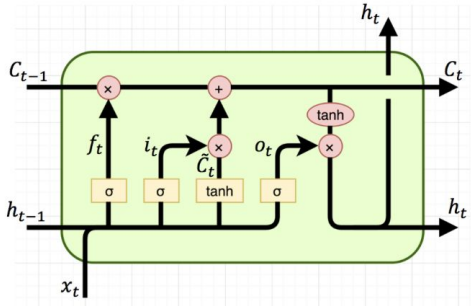

- increasing complexity
- increasing train time
- slower convergence

Source:
https://medium.com/@saurabh.rathor092/simple-rnn-vs-gru-vs-lstm-difference-lies-in-more-flexible-control-5f33e07b1e57

Simple RNN

no pretrained word vectors

GRU + LSTM

LexVec

	f1-score	f1-score
Computer Science	0.823	0.811
Physics	0.874	0.879
Mathematics	0.804	0.806
Statistics	0.720	0.785
Quantitative Biology	0.000	0.115
Quantitative Finance	0.000	0.043
micro avg	0.798	0.807
macro avg	0.537	0.573
weighted avg	0.784	0.798
samples avg	0.805	0.801

Bidirectional Encoder Representations from Transformers

BERT

1) 12 layers of encoder and decoder
2) attention mechanism

Layer (type)	Output Shape	Param #	Connected to
Input-Token (InputLayer)	[(None, 128)]	0	
Input-Segment (InputLayer)	[(None, 128)]	0	
Embedding-Token (TokenEmbedding	[(None, 128, 768), (	23440896	Input-Token[0][0]
Embedding-Segment (Embedding)	(None, 128, 768)	1536	Input-Segment[0][0]
Embedding-Token-Segment (Add)	(None, 128, 768)	0	Embedding-Token[0][0] Embedding-Segment[0][0]
Embedding-Position (PositionEmb	(None, 128, 768)	98304	Embedding-Token-Segment[0][0]
Embedding-Dropout (Dropout)	(None, 128, 768)	0	Embedding-Position[0][0]
Embedding-Norm (LayerNormalizat	(None, 128, 768)	1536	Embedding-Dropout[0][0]
Encoder-1-MultiHeadSelfAttentio	(None, 128, 768)	2362368	Embedding-Norm[0][0]
Encoder-1-MultiHeadSelfAttentio	(None, 128, 768)	0	Encoder-1-MultiHeadSelfAttention[
Encoder-1-MultiHeadSelfAttentio	(None, 128, 768)	0	Embedding-Norm[0][0] Encoder-1-MultiHeadSelfAttention-
Encoder-1-MultiHeadSelfAttentio	(None, 128, 768)	1536	Encoder-1-MultiHeadSelfAttention-
Encoder-1-FeedForward (FeedForw	(None, 128, 768)	4722432	Encoder-1-MultiHeadSelfAttention-
Encoder-1-FeedForward-Dropout (	(None, 128, 768)	0	Encoder-1-FeedForward[0][0]
Encoder-1-FeedForward-Add (Add)	(None, 128, 768)	0	Encoder-1-MultiHeadSelfAttention- Encoder-1-FeedForward-Dropout[0][
Encoder-1-FeedForward-Norm (Lay	(None, 128, 768)	1536	Encoder-1-FeedForward-Add[0][0]

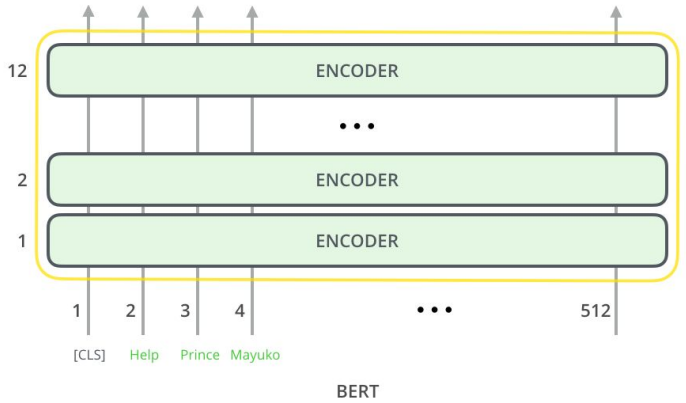

Size Matters

	BERT	**distilBERT**
max sequence	128	no limitation (300)
parameters	109,191,942	66,958,086
output model size	1.3 GB	250 MB

DistilBERT

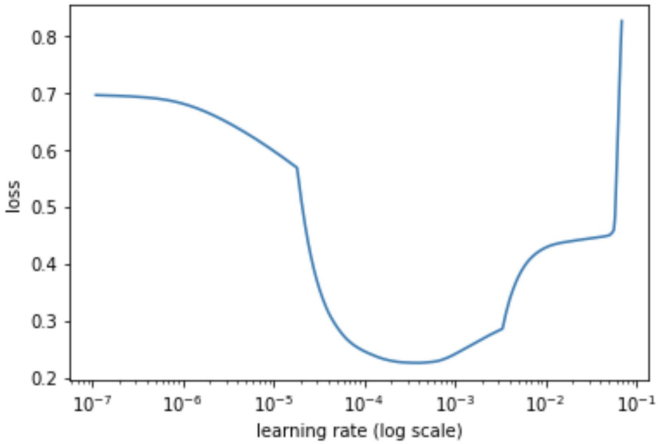
loss
learning rate (log scale)

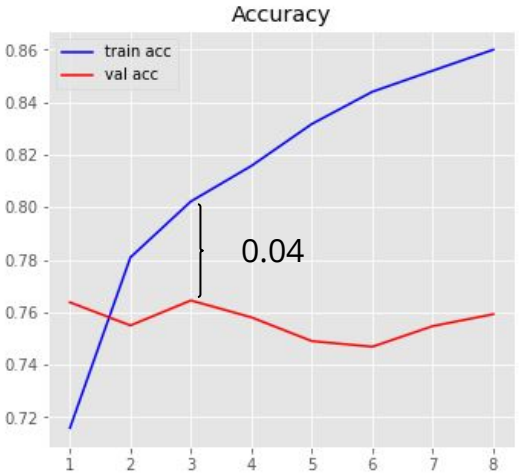
Accuracy
train acc
val acc
0.04

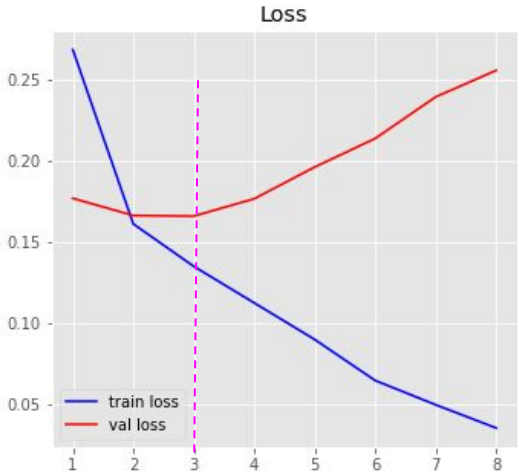
Loss
train loss
val loss

Final Model - distilBERT

	precision	recall	f1-score	support
Computer Science	0.819	0.884	0.850	1692
Physics	0.911	0.874	0.892	1226
Mathematics	0.852	0.774	0.811	1150
Statistics	0.776	0.809	0.792	1069
Quantitative Biology	0.602	0.557	0.579	122
Quantitative Finance	0.889	0.711	0.790	45
micro avg	0.833	0.834	0.833	5304
macro avg	0.808	0.768	0.786	5304
weighted avg	0.835	0.834	0.833	5304
samples avg	0.860	0.869	0.845	5304

hamming loss : 0.07024235200635677

Summary of Results

library	model	type	micro F1	hamming loss	classes with nil prediction	lowest recall
sklearn	logistic regression	one vs rest	0.799	0.0709	0	0.041
sklearn	support vector machine	one vs rest	0.811	0.0756	0	0.016
keras	simple neural network without word vectors	multilabel	0.798	0.0825	2	0.000
keras	simple RNN with LexVec	multilabel	0.677	0.1205	3	0.000
keras	LSTM and GRU with LexVec	multilabel	0.807	0.0776	0	0.066
ktrain	biGRU	multilabel	0.82	0.0733	0	0.356
ktrain	BERT	multilabel	0.825	0.0729	0	0.352
ktrain	distilBERT	multilabel	0.833	0.0702	0	0.557

Is that the ground truth?

“Efficient and consistent inference of ancestral sequences in an evolutionary model with insertions and deletions under dense taxon sampling. In evolutionary biology, the speciation history of living organisms is represented graphically by a phylogeny, that is, a rooted tree whose leaves correspond to current species and branchings indicate past speciation events. Phylogenies are commonly estimated from molecular sequences, such as DNA sequences, collected from the species of interest.”

hamming loss = 4/6

	Com Sci	Physics	Math	Statistics	Quant. Biology	Quant. Finance
ground truth	**1**	**0**	**1**	**1**	**0**	**0**
prediction	**0**	**0**	**0**	**0**	**1**	**0**

Is that the ground truth?

“Cyclic Dominance in the Spatial Coevolutionary Optional Prisoner's Dilemma Game This paper studies scenarios of cyclic dominance in a coevolutionary spatial model in which game strategies and links between agents adaptively evolve over time. The Optional Prisoner's Dilemma (OPD) game is employed. The OPD is an extended version of the traditional Prisoner's Dilemma where players have a third option to abstain from playing the game. We adopt an agent-based simulation approach and use Monte Carlo methods to perform the OPD with coevolutionary rules. The necessary conditions to break the scenarios of cyclic dominance are also investigated. This work highlights that cyclic dominance is essential in the sustenance of biodiversity.”

hamming loss = 4/6

	Com Sci	Physics	Math	Statistics	Quant. Biology	Quant. Finance
ground truth	**1**	**1**	**1**	**0**	**0**	**0**
prediction	**0**	**0**	**0**	**0**	**1**	**0**

Suggested Further Studies

1) Tune other hyperparameters such as optimizer and learning rates.

2) Summarize text using BART or T5 to half the length and repeat training using BERT.

3) Relabel or remove the data with obviously wrong ground truths.

Thank you!

Teo Zhan Rui
teozhanrui@gmail.com

www.linkedin.com/in/teozhanrui

https://github.com/teozhanrui